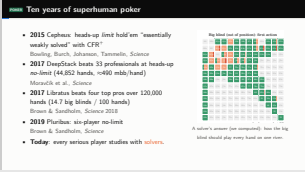


Timing plan (60 min). Hook & rules 6 · Games, equilibria, sequence form 12 · Regret minimisation and CFR 14 · What solutions look like 6 · Our method 14 · Results, limits, wrap-up 8.
If running late, skip: "Kuhn poker", "Modern CFR = numerical tricks" (one sentence), "Attempt 1: Anderson" (one sentence), backup slides.

Opening line: "Ten years ago the best poker players in the world lost to programs whose core is a few lines of arithmetic repeated a few billion times. This talk is about that arithmetic, read as numerical analysis — and about what a numerical analyst can add to it."

The slides alternate between POKER (green tag) and MATH (blue tag); orange NEW marks our own contribution. Everything labelled "we computed" was produced by the Python solver that accompanies the talk (repository `poker-solver`); everything else is quoted from the cited papers.

POKER Ten years of superhuman poker

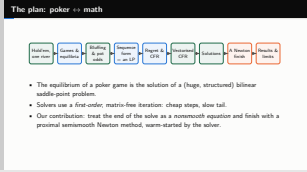


The common core of all these programs is an iterative method — counterfactual regret minimisation — applied to a gigantic saddle-point problem. Libratus and DeepStack add *abstraction* (shrinking the game) and *subgame solving* (re-solving the part of the game you are actually in), but the inner loop is still a CFR-type method.

"Solvers" (PioSolver, GTO+, GTO Wizard, ...) are what human players use: you describe one spot — the board, the two ranges, the allowed bet sizes — and the program computes an equilibrium of that subgame. The picture on the right is exactly such an output, produced by our own solver for a river spot we will meet again.

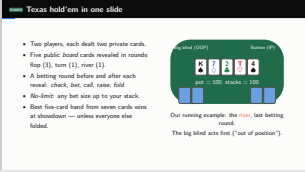
mbb/hand = thousandths of a big blind per hand.

The plan: poker ↔ math



Say what the audience will get: (1) enough poker to read a solver output, (2) the mathematics of why CFR converges and why it is fast in practice, (3) a new method — from the numerical-optimisation toolbox, not the game-theory one — with honest results: where it wins by orders of magnitude, and where it does not (yet).

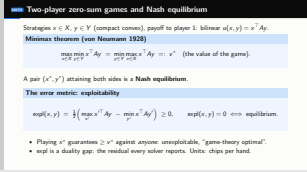
POKER Texas hold'em in one slide



Keep this light — most of the audience has played or watched poker. The key structural fact: the board is *public*, the two hole cards are *private*. Strategy must therefore be a function of (public history, your own cards), and you do not know the opponent's cards — only a *probability distribution* over them, which poker players call a **range**.

The river spot is fixed for the whole talk: board K♠ 7♢ 2♣ T♡ 4♠, pot 100, both players have 100 behind. Big blind may check or bet 33%, 75% of pot or all-in; facing a bet you may fold, call or go all-in.

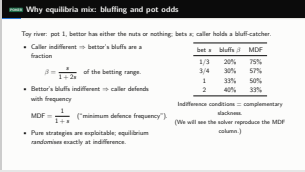
MATH Two-player zero-sum games and Nash equilibrium



For an applied-math audience: this is a convex–concave saddle-point problem with a bilinear coupling; exploitability is the primal–dual gap. Everything later is a method for driving this gap to zero.

Practical solvers stop at an exploitability of a fraction of a percent of the pot; we will push to 10^{-13}. Our numbers are in chips with pot $= 100$, so 10^{-2} chips is 10^{-4} of the pot.

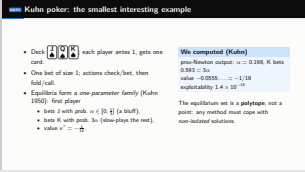
POKER Why equilibria mix: bluffing and pot odds



Derivation (say it, don't write it). Caller with a bluff-catcher: calling wins $1+s$ against a bluff and loses s against value; indifferent when $\beta(1+s) = (1-\beta)s$, i.e. $\beta = s/(1+2s)$. Bettor with air: betting wins the pot 1 when the caller folds and loses s when called; checking wins nothing; indifferent when $(1-c)\cdot 1 = c\,s$, i.e. $c = 1/(1+s)$.

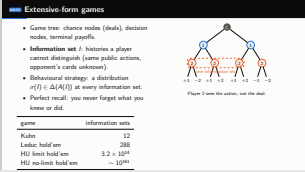
The mathematical point: at an equilibrium, every action played with positive probability is a best response — this is complementary slackness in the LP we are about to write down. Mixing happens exactly where reduced costs vanish.

MATH Kuhn poker: the smallest interesting example



Kuhn poker is the poker analogue of a 2×2 test matrix. The one-parameter family matters for numerical analysis: the solution set of the equilibrium problem is a polyhedron, generally not a point, so Jacobians at a solution are singular and textbook Newton theory does not apply as-is. Our method handles this by regularisation (proximal point) — which also picks one particular member of the family (here $\alpha \approx 0.198$).

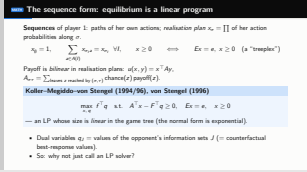
MATH Extensive-form games



Leduc hold'em (Southey et al. 2005) is the standard research benchmark: a six-card deck (J, Q, K in two suits), one private card each, one public card, two betting rounds with bets 2 and 4 and at most two raises per round. 288 information sets.

The size column is the reason the subject exists: one cannot write down HUNL. Real systems solve an *abstraction* (similar hands and bet sizes merged) and then re-solve subgames like our river spot during play.

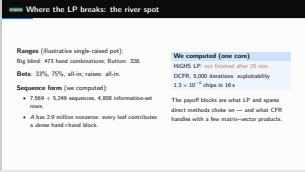
MATH The sequence form: equilibrium is a linear program



This is the key modelling step. Mixed strategies over pure strategies of the normal form would have exponentially many variables; realisation plans have one variable per sequence, and perfect recall makes the payoff bilinear in them.

E and F are very sparse (one row per information set, entries ± 1) and encode flow conservation on a tree. A is where the poker lives. Remember the dual variables: they are *values* of information sets, i.e. the same quantities CFR computes every iteration — this is how we will warm-start the duals later.

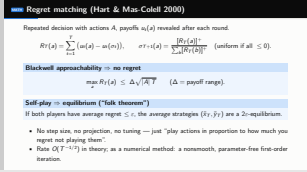
POKER Where the LP breaks: the river spot



This is the honest motivation for first-order methods: for subgame-sized problems with full ranges, the constraint matrix is not "sparse" in any useful sense, because each terminal couples every hand of one player to every compatible hand of the other. The LP solver run (HiGHS via SciPy, default method) shared the machine with another job; we stopped it after 25 minutes.

On the smaller test spots the LP is perfectly fine (Leduc: 0.04 s, small river: 1.6 s) — we will show those numbers too.

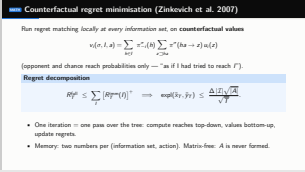
MATH Regret matching (Hart & Mas-Colell 2000)



Proof idea of the bound: the regret vector R_T satisfies $\|[R_T]^+\|^2 \le \|[R_{T-1}]^+\|^2 + \Delta^2|A|$ because the cross term $\langle [R_{T-1}]^+, r_T\rangle$ vanishes by the choice of σ_T — the instantaneous regret is orthogonal to the current positive regrets. Sum and take square roots.

The folk theorem: average regret ε for both players means $\max_{x'} x'^{\top} A\bar{y} - \min_{y'} \bar{x}^{\top} A y' \le 2\varepsilon$, because the average of the realised payoffs is sandwiched.

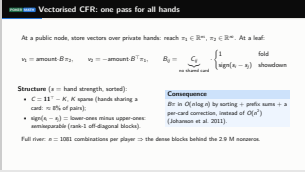
MATH Counterfactual regret minimisation (Zinkevich et al. 2007)



The decomposition is the whole trick: a global regret over an exponential strategy space is bounded by a sum of small, local regrets, each minimised by plain regret matching. This is why CFR scales: cost per iteration is linear in the tree, storage is linear in the number of information-set actions.

In linear-algebra language one CFR iteration computes Ay and $A^\top x$ (with x, y the current realisation plans) by tree traversals — the same two products a first-order saddle-point method needs.

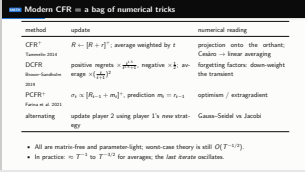
POKER MATH Vectorised CFR: one pass for all hands



This is the poker-specific numerical linear algebra. Each terminal block is "low rank plus sparse": the fold block is rank one minus a sparse card-removal correction; the showdown block, after sorting hands by strength, is a sign matrix (a semiseparable matrix — think of the inverse of a bidiagonal matrix) times the compatibility mask. Products with it are prefix sums.

For our sizes (a few hundred to 1081 hands) a dense BLAS product is already fast, so our code simply uses dense B; but this structure is what makes river solvers fast in practice, and it is the structure a scalable Newton method would have to exploit (end of talk).

MATH Modern CFR = a bag of numerical tricks

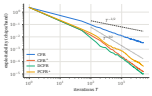


CFR$^+$ is what solved limit hold'em. DCFR with $(\alpha, \beta, \gamma) = (1.5, 0, 2)$ is the default in many solvers. PCFR$^+$ connects regret matching to optimistic mirror descent (Farina, Kroer, Sandholm, AAAI 2021).

The empirical rates are much better than the $T^{-1/2}$ worst case — which is one of the open theoretical questions in the area — but they are still sublinear, and that is the gap our method addresses.

MATH We computed: the baselines on Leduc hold'em

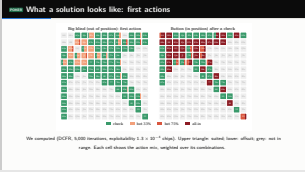
MATH We computed: the baselines on Leduc hold'em



- Our range-vectorised implementation (alternating updates), exploitability of the average strategy.
- 5,000 iterations, 14×: DCFR reaches 1.0×10^{-6}.
- 50,000 iterations, 130×: CFR$^+$ 9.2×10^{-7}, DCFR 1.6×10^{-6}.
- Each extra digit costs $\approx 10\times$ more iterations.

Exploitability is in chips per hand (Leduc antes are 1 chip). The ordering CFR $\ll$ PCFR$^+$ $<$ CFR$^+$ $\approx$ DCFR is the usual one in the literature for Leduc. Note the tail: to go from 10^{-5} to 10^{-6} took ten times the work.

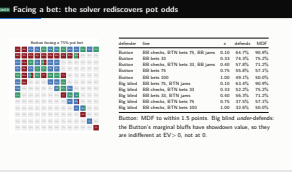
POKER What a solution looks like: first actions



Read the grid like a solver user. Big blind (left): checks most of the range, uses a small 33% bet with medium hands (Q-x, J-x type hands) — a "block bet" — and almost never bets big. Button after a check (right): *polarised* — jams the very strong hands (sets, KQ+, AK) and a set of bluffs chosen by card removal (A5s, A6s, 65s, 86s . . .), checks back the medium hands that win at showdown.

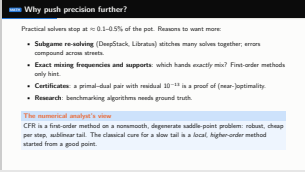
Aggregate frequencies (we computed): big blind checks 82%, bets 33% pot 16%; Button after a check jams 44%, checks 50%.

POKER Facing a bet: the solver rediscovers pot odds



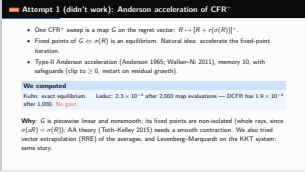
This is a nice consistency check of the whole pipeline. When the big blind bets, the Button defends 74.3% vs 75.2%, 55.8% vs 57.1%, 49.1% vs 50.0% — the toy-game formula is almost exact. When the Button bets after a check, the big blind folds much more than the MDF: we checked that the Button's weakest hands mix between checking and jamming, so the indifference that pins the defence frequency is "EV(jam) = EV(check) > 0", not "EV(jam) $= 0$". The toy model's assumption (bluffs have no showdown value) fails, and the solver tells you by how much.

MATH Why push precision further?



Make the case without overselling: for a human studying a spot, 0.3% of the pot is plenty. Precision matters for machines (re-solving, evaluation) and for science (knowing the exact equilibrium, e.g. the exact bluffing frequencies, or whether a hand mixes at all). The analogy for this audience: CFR is like a stationary iterative method — fantastic for the first digits, hopeless for the last ones — and the obvious idea is to switch to a Newton-type method near the solution.

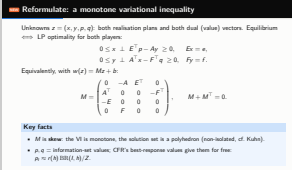
NEW Attempt 1 (didn't work): Anderson acceleration of CFR$^+$



Include this slide: audiences of numerical analysts appreciate negative results, and it motivates the reformulation. The lesson we drew: do not accelerate the regret dynamics; go back to the optimality conditions themselves and fix their two defects — non-uniqueness (handled by a proximal term) and nonsmoothness (handled by a semismooth Newton method on a smoothed-at-the-kink reformulation).

Also tried (we computed, not shown): reduced-rank extrapolation of the sequence of DCFR averages gave no improvement; plain semismooth Levenberg–Marquardt on the min-function KKT system stalled on Leduc (degenerate complementarity).

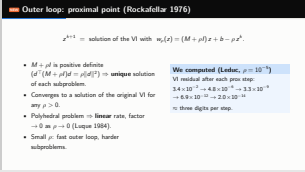
NEW Reformulate: a monotone variational inequality



This is a mixed linear complementarity problem. The reduced cost $E^\top p - Ay$ of a sequence is its counterfactual *regret*, scaled by reach: complementarity says "play only actions with zero regret". So CFR and this system talk about the same quantities — CFR iterates give us x, y (average strategies in sequence form) and, via one best-response pass, p, q.

Skewness: the x–y block is the zero-sum coupling; the (x, p) and (y, q) blocks are constraint/multiplier pairs, arranged with signs so that $M^\top = -M$.

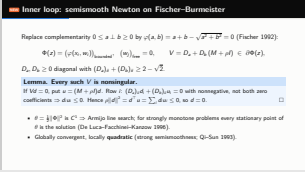
NEW Outer loop: proximal point (Rockafellar 1976)



The proximal term fixes non-uniqueness: each subproblem has exactly one solution, so Newton's method has a well-defined, nonsingular target. For polyhedral monotone problems the inverse operator is upper Lipschitz (Robinson 1981), which gives linear convergence of the proximal point method with a factor that tends to zero as $\rho \to 0$ (Luque 1984; Rockafellar 1976 for the general theory).

Table from our runs: at $\rho = 10^{-3}$ Leduc needed 19 prox steps, at 10^{-4}: 6, at 10^{-5}: 4, at 10^{-6}: 3. We use $\rho = 10^{-5}$.

NEW Inner loop: semismooth Newton on Fischer–Burmeister


Inner loop: semismooth Newton on Fischer–Burmeister

Replace complementarity $0 \le a \perp b \ge 0$ by $\varphi(a,b) = a + b - \sqrt{a^2 + b^2} = 0$ (Fischer 1992):

$$\Phi(x) = \big(\varphi(v_i, w_i)\big)_{i \in \text{unmatched}}, \quad (w)_{\text{free}} = 0, \qquad V = D_v + D_b(M + \rho I) \in \partial\Phi(x),$$

$D_v, D_b \ge 0$ diagonal with $(D_v)_{ii} + (D_b)_{ii} \ge 2 - \sqrt{2}$.

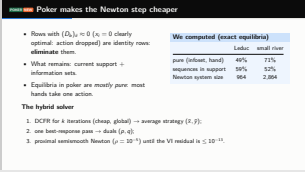
Lemma. Every such V is nonsingular.
If $Vd = 0$, put $u = (M + \rho I)d$. Row i: $(D_v)_{ii} d_i + (D_b)_{ii} u_i = 0$ with nonnegative, not both zero coefficients $\Rightarrow d_i u_i \le 0$. Hence $\rho \|d\|^2 = d^\top u = \sum_i d_i u_i \le 0$, so $d = 0$. □

- $\theta = \tfrac{1}{2}\|\Phi\|^2$ is $C^1 \Rightarrow$ Armijo line search; for strongly monotone problems every stationary point of θ is the solution (De Luca–Facchinei–Kanzow 1996).
- Globally convergent, locally **quadratic** (strong semismoothness; Qi–Sun 1993).

Why Fischer–Burmeister rather than $\min(a, b)$: the min-function gives a Newton method (the primal–dual active-set method) whose merit function is nonsmooth, and on our river spot the line search on it stalled with steps of 10^{-4}, while full steps diverged. FB has a continuously differentiable merit function, so standard globalisation works.

The lemma is the reason for the proximal term: without ρ, $d^\top M d = 0$ and the argument fails — indeed Newton matrices of the unregularised problem are singular at degenerate solutions, which is exactly what we observed with Levenberg–Marquardt.

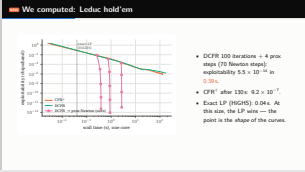
POKER NEW Poker makes the Newton step cheaper



Mixing happens only at indifference, which is non-generic: in the exact Leduc equilibrium about half of the reachable (information set, hand) pairs play a single action; on the small river spot 71%. The reduced Newton system keeps only sequences that are still "undecided".

Honest remark: the reduction is moderate on these spots (sizes 65–90% of the full system during the solve) because the Fischer–Burmeister derivative only becomes exactly negligible once the iterate is close; a cleverer active-set prediction from the CFR iterate is an obvious improvement.

NEW We computed: Leduc hold'em

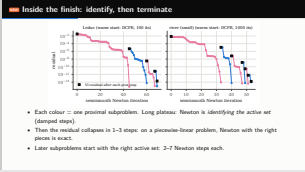


Each pink curve starts where the DCFR warm start stopped (100, 300 or 1000 iterations) and then drops vertically: three digits per proximal step. Starting from 100 DCFR iterations is enough; starting later only costs DCFR time.

The game value we obtain, $-0.085606424078\ldots$ chips for the first player, agrees with the LP to all printed digits.

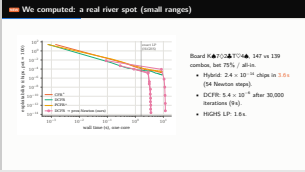
Be upfront about the LP: Leduc is tiny (964 unknowns); a simplex method finishes in 40 ms. The question is what happens when the LP does not fit — next two slides.

NEW Inside the finish: identify, then terminate



This is the characteristic signature of semismooth Newton on complementarity problems: a globalisation phase in which the line search accepts partial steps while the active set changes, followed by finite termination (up to rounding) once the active set is right. Note that the warm start matters mostly for the length of the plateau.

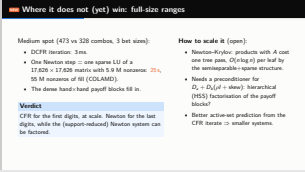
NEW We computed: a real river spot (small ranges)



Same picture on a genuine no-limit river subgame with card removal and a non-trivial betting tree. The hybrid reaches machine precision about 10^8 times more accurately than DCFR in less than half the time; DCFR's tail would need millions of iterations for the same accuracy. Here the LP is still competitive.

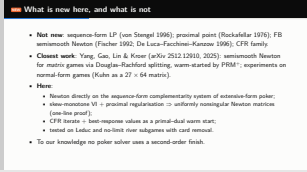
Exploitability in chips with pot 100: 10^{-6} chips is 10^{-8} of the pot.

NEW Where it does not (yet) win: full-size ranges



State the limitation plainly. With a direct sparse solver the Newton step costs as much as roughly 8,000 DCFR iterations on this spot, because the factorisation sees the payoff blocks as dense. A matrix-free Krylov method would have per-iteration cost equal to a CFR iteration, but the shifted skew system with $\rho = 10^{-5}$ is badly conditioned (eigenvalues $\rho + i\lambda$), so it needs a structured preconditioner. The structure is there — semiseparable plus sparse — but building the solver is future work.

NEW What is new here, and what is not



Be precise about novelty. The general idea “first-order warm start, then semismooth Newton” for zero-sum games was proposed for matrix games by Yang, Gao, Lin and Kroer in December 2025; their Newton system lives on the normal form and they do not treat extensive-form structure. Our contribution is the extensive-form, sequence-form version with a different regularisation (proximal point on the skew VI), a CFR-based dual warm start, and the poker experiments. Commercial and research poker solvers we know of (PioSolver, GTO+, Libratus, DeepStack) are first-order (CFR variants).

Takeaways

1. A poker equilibrium is the solution of a structured bilinear saddle-point problem; the sequence form makes it an LP of size linear in the tree.
2. CFR = regret matching at every information set: a matrix-free, nonsmooth first-order method with a provable $O(T^{-1/2})$ rate and a much better practical one — plus numerical tricks (projection, discounting, optimism, Gauss–Seidel).
3. The solver's output is readable poker: polarised betting, block bets, pot-odds defence (to within 1.5%).
4. Read the equilibrium conditions as a **skew-monotone VI**: proximal point + Fischer–Burmeister semismooth Newton turns CFR's slow tail into a few linear solves and reaches 10^{-13}.
5. Scaling the finish is a structured numerical linear algebra problem (semiseparable + sparse blocks).

Code: `github.com/deanguambers/Poker-Solver` — solver, benchmarks, figures, these slides.

End with the invitation: the open problem on the previous slide is a natural one for this club — a preconditioner for shifted skew systems whose off-diagonal blocks are semiseparable-plus-sparse.

Backup only.

Backup: implementation summary

`games.py`	Kuhn, Leduc, no-limit river builder (7-card evaluator, card removal, range notation)
`cfr.py`	range-vectorised public-tree CFR / CFR$^+$ / DCFR / PCFR$^+$, alternating updates, best response, exploitability
`seqform.py`	sequence form (E, F, A), dual estimates from best responses, prox semismooth Newton (FB, reduced)
`lp.py`	reference LP via HiGHS
`aa.py`, `extrap.py`	the negative results (Anderson, RRE)
`bench_*.py`, `fig_*.py`	every number and figure in these slides

- Exploitability: $\frac{1}{2}(\mathrm{BR}_1 + \mathrm{BR}_2)$ by exact best-response passes, normalised by the deal weight Z.
- Sequence form checked against the LP: Leduc value -0.085604240780 (both).
- Dual warm start: $p_i \propto r_i(h)\,\mathrm{BR}_i(I,h)/Z$, $q_i \propto -r_i(h)\,\mathrm{BR}_i(J,h)/Z$.

Backup only.

Backup: effect of the proximal parameter

We computed (Leduc, warm start DCFR 200 iterations, final exploitability $\leq 10^{-11}$ in all cases):

ρ	10^{-3}	10^{-4}	10^{-5}	10^{-6}
prox steps	19	6	4	3

Continuation ($\rho_k \downarrow$) did not reduce the total number of Newton steps on the small river spot (54 fixed vs 63–96).